




Heterogeneous Parallel Programming



Lecture 4.1 Parallel Computation Patterns Reduction

Wen-mei Hwu - University of Illinois at Urbana-Champaign

Objective

- To learn the parallel reduction pattern
 - An important class of parallel computation
 - Work efficiency analysis
 - Resource efficiency analysis

Partition and Summarize

- A commonly used strategy for processing large input data sets
 - There is no required order of processing elements in a data set (associative and commutative)
 - Partition the data set into smaller chunks
 - Have each thread to process a chunk
 - Use a reduction tree to summarize the results from each chunk into the final answer
- Google and Hadoop MapReduce frameworks support this strategy
- We will focus on the reduction tree step for now.

Reduction enables other techniques.

- Reduction is also needed to clean up after some commonly used parallelizing transformations
- Privatization
 - Multiple threads write into an output location
 - Replicate the output location so that each thread has a private output location
 - Use a reduction tree to combine the values of private locations into the original output location

What is a reduction computation?

- Summarize a set of input values into one value using a “reduction operation”
 - Max
 - Min
 - Sum
 - Product
- Often with user defined reduction operation function as long as the operation
 - Is associative and commutative
 - Has a well-defined identity value (e.g., 0 for sum)

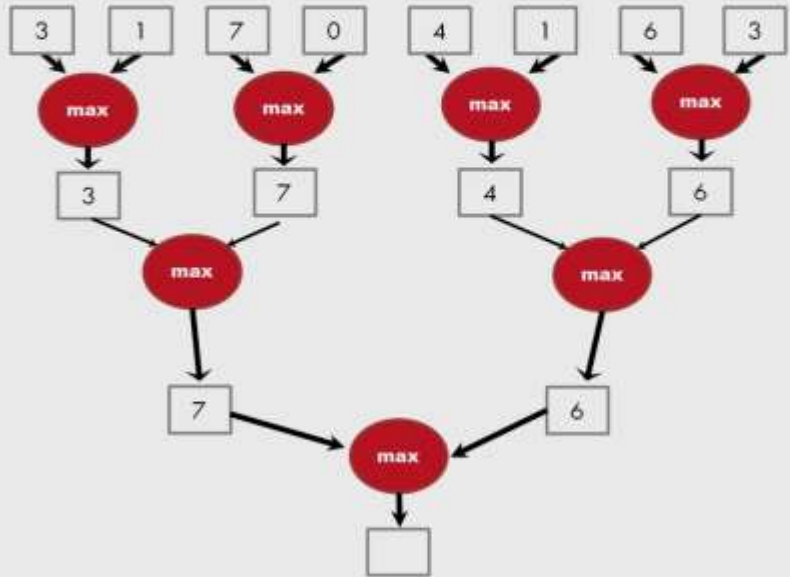
An example of “collective operation”

An Efficient Sequential Reduction $O(N)$


- Initialize the result as an identity value for the reduction operation
 - Smallest possible value for max reduction
 - Largest possible value for min reduction
 - 0 for sum reduction
 - 1 for product reduction
- Iterate through the input and perform the reduction operation between the result value and the current input value
 - N reduction operations performed for N input values



A parallel reduction tree algorithm performs $N-1$ Operations in $\log(N)$ steps




A tournament is a reduction tree with “max” operation



A Quick Analysis

- For N input values, the reduction tree performs
 - $(1/2)N + (1/4)N + (1/8)N + \dots (1)N = (1 - (1/N))N = N-1$ operations
 - In $\log(N)$ steps - 1,000,000 input values take 20 steps
 - Assuming that we have enough execution resources
 - Average Parallelism $(N-1)/\log(N)$
 - For $N = 1,000,000$, average parallelism is 50,000
 - However, peak resource requirement is 500,000!
 - This is not resource efficient.
- This is a work-efficient parallel algorithm
 - The amount of work done is comparable to sequential
 - Many parallel algorithms are not work efficient



To learn more, read
Section 6.1