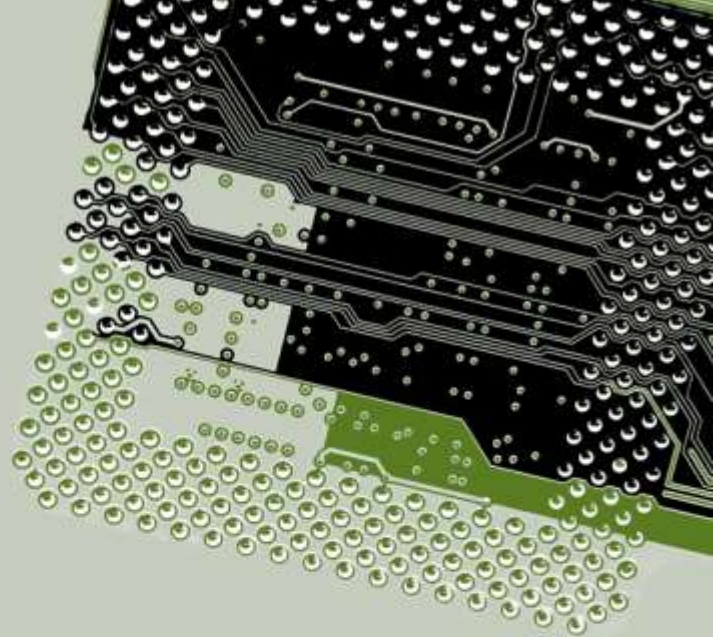




Heterogeneous Parallel Programming



Lecture 4.2

Parallel Computation Patterns

A Basic Reduction Kernel

Wen-mei Hwu - University of Illinois at Urbana-Champaign

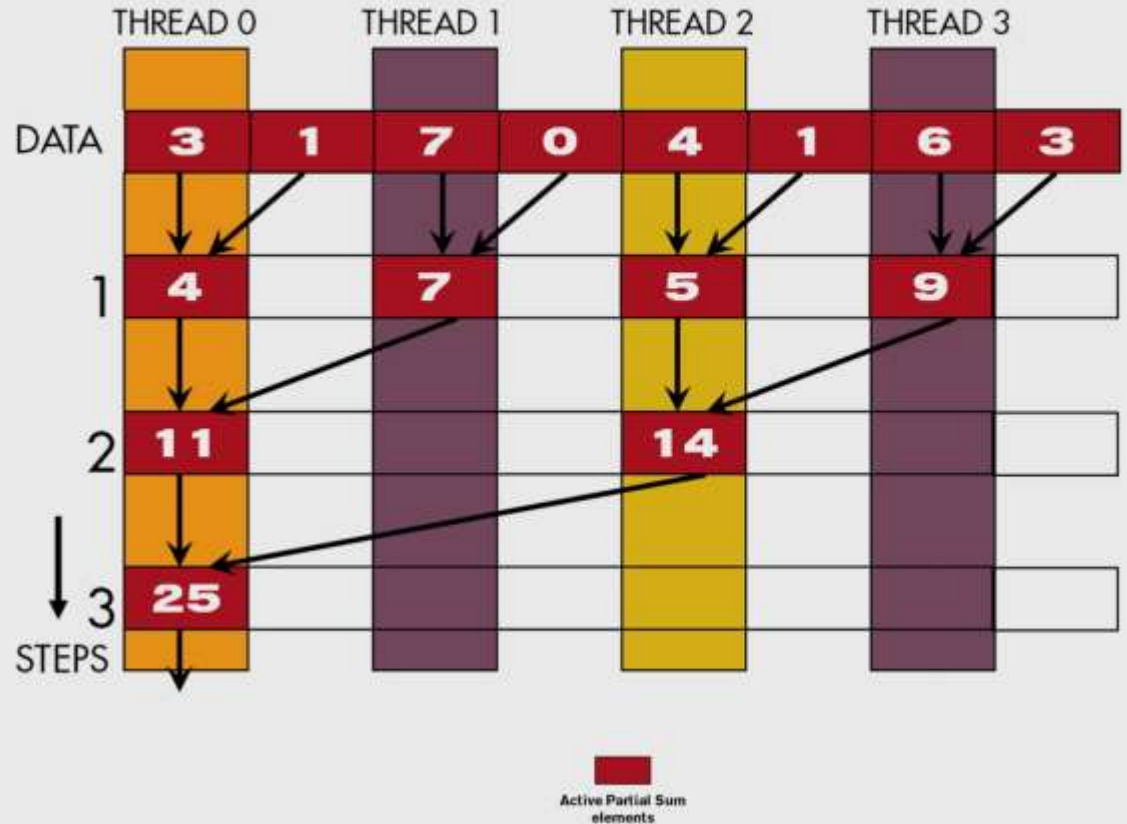
Objective

- To learn to write a basic reduction kernel
 - Thread to data index mapping
 - Turning off threads
 - Control divergence

Parallel Sum Reduction

- Parallel implementation:
 - Recursively halve # of threads, add two values per thread in each step
 - Takes $\log(n)$ steps for n elements, requires $n/2$ threads
- Assume an in-place reduction using shared memory
 - The original vector is in device global memory
 - The shared memory is used to hold a partial sum vector
 - Each step brings the partial sum vector closer to the sum
 - The final sum will be in element 0
 - Reduces global memory traffic due to partial sum values
 - Thread block size limits n to be less than or equal to 2,048

A Parallel Sum Reduction Example



A Naive Thread Index to Data Mapping

- Each thread is responsible of an even-index location of the partial sum vector
- After each step, half of the threads are no longer needed
- One of the inputs is always from the location of responsibility
- In each step, one of the inputs comes from an increasing distance away

A Simple Thread Block Design

- Each thread block takes $2 * \text{BlockDim.x}$ input elements
- Each thread loads 2 elements into shared memory

```
__shared__ float partialSum[2*BLOCK_SIZE];  
  
unsigned int t = threadIdx.x;  
unsigned int start = 2*blockIdx.x*blockDim.x;  
partialSum[t] = input[start + t];  
partialSum[blockDim+t] = input[start + blockDim.x+t];
```

The Reduction Steps

```
for (unsigned int stride = 1;
     stride <= blockDim.x;  stride *= 2)
{
    __syncthreads();
    if (t % stride == 0)
        partialSum[2*t] +=
        partialSum[2*t+stride];
}
```

Why do we need syncthreads()?

Barrier Synchronization

- `__syncthreads()` are needed to ensure that all elements of each version of partial sums have been generated before we proceed to the next step

Back to the Global Picture

- At the end of the kernel, Thread 0 in each thread block writes the sum of the thread block in `partialSum[0]` into a vector indexed by the `blockIdx.x`
- There can be a large number of such sums if the original vector is very large
 - The host code may iterate and launch another kernel
- If there are only a small number of sums, the host can simply transfer the data back and add them together.



To learn more, read [Section 6.1](#)

