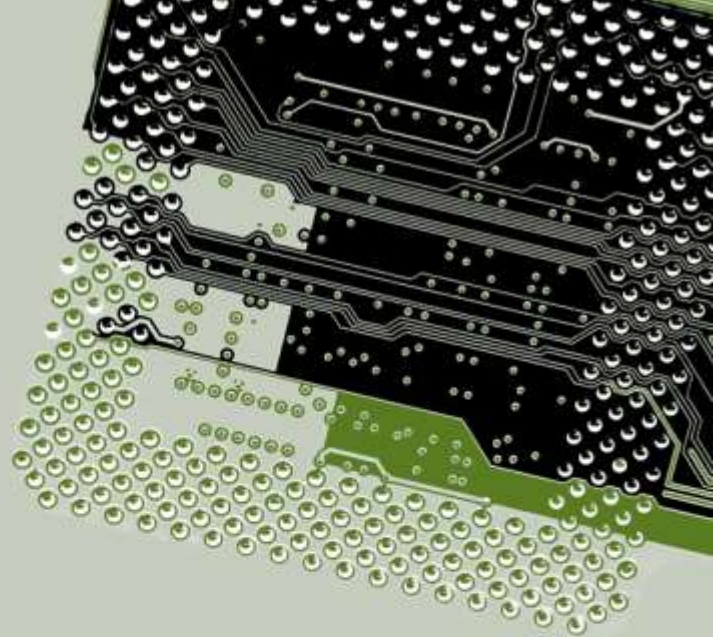




Heterogeneous Parallel Programming



Lecture 4.6

Parallel Computation Patterns

A Work-Efficient Parallel Scan Kernel

Wen-mei Hwu - University of Illinois at Urbana-Champaign

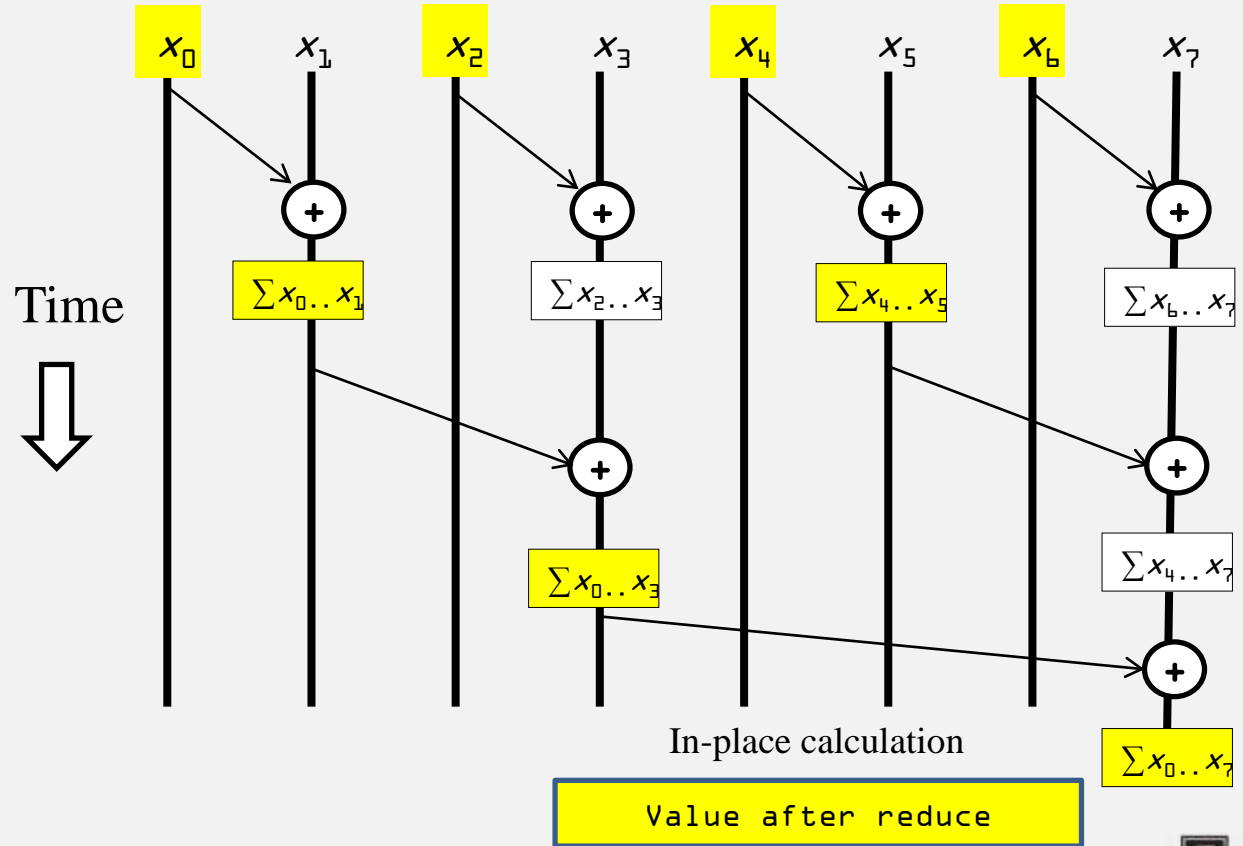
Objective

- To learn to write a work-efficient scan kernel
 - Two-phased balanced tree traversal
 - Aggressive reuse of computation results
 - Reducing control divergence with more complex thread index to data index mapping

Improving Efficiency

- *Balanced Trees*
 - Form a balanced binary tree on the input data and sweep it to and from the root
 - Tree is not an actual data structure, but a concept to determine what each thread does at each step
- For scan:
 - Traverse down from leaves to root building partial sums at internal nodes in the tree
 - Root holds sum of all leaves
 - Traverse back up the tree building the output from the partial sums

Parallel Scan - Reduction Phase



Reduction Phase Kernel Code

```
// XY[2*BLOCK_SIZE] is in shared memory

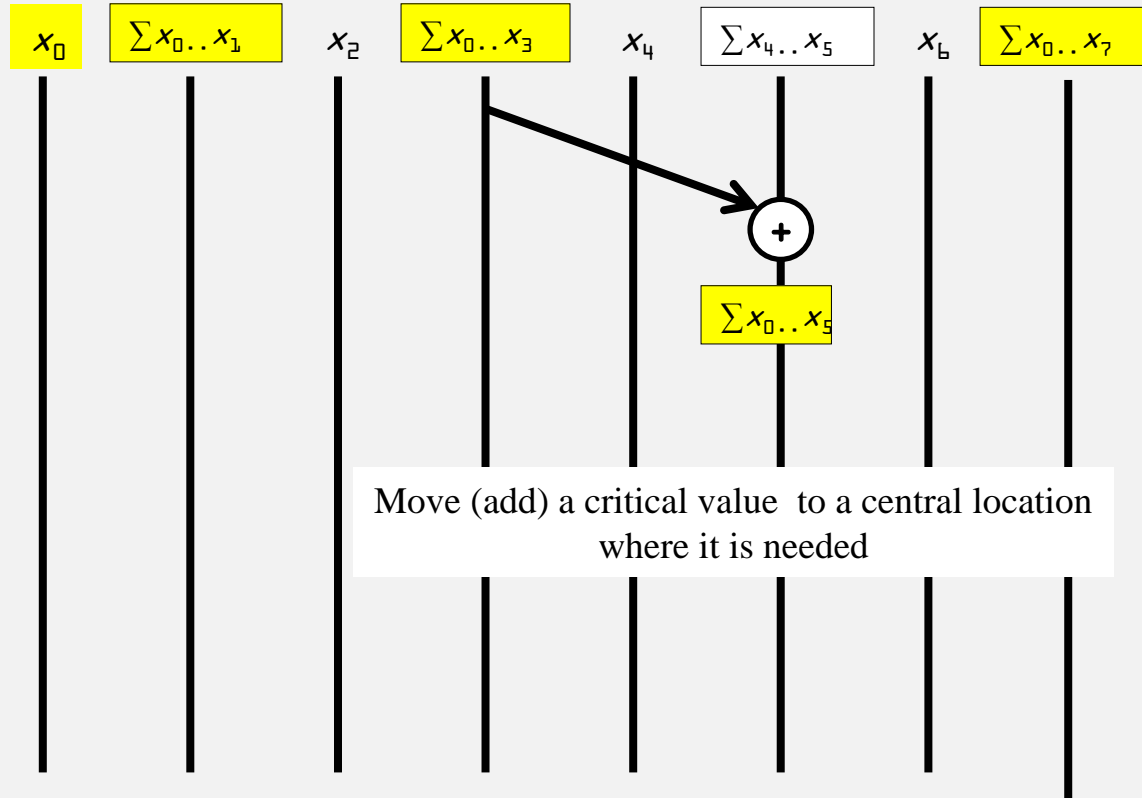
for (int stride = 1; stride <= BLOCK_SIZE;
     stride *= 2) {
    int index = (threadIdx.x+1)*stride*2 - 1;
    if(index < 2*BLOCK_SIZE)
        XY[index] += XY[index-stride];
    __syncthreads();
}
```

$\text{threadIdx.x}+1 = 1, 2, 3, 4, \dots$

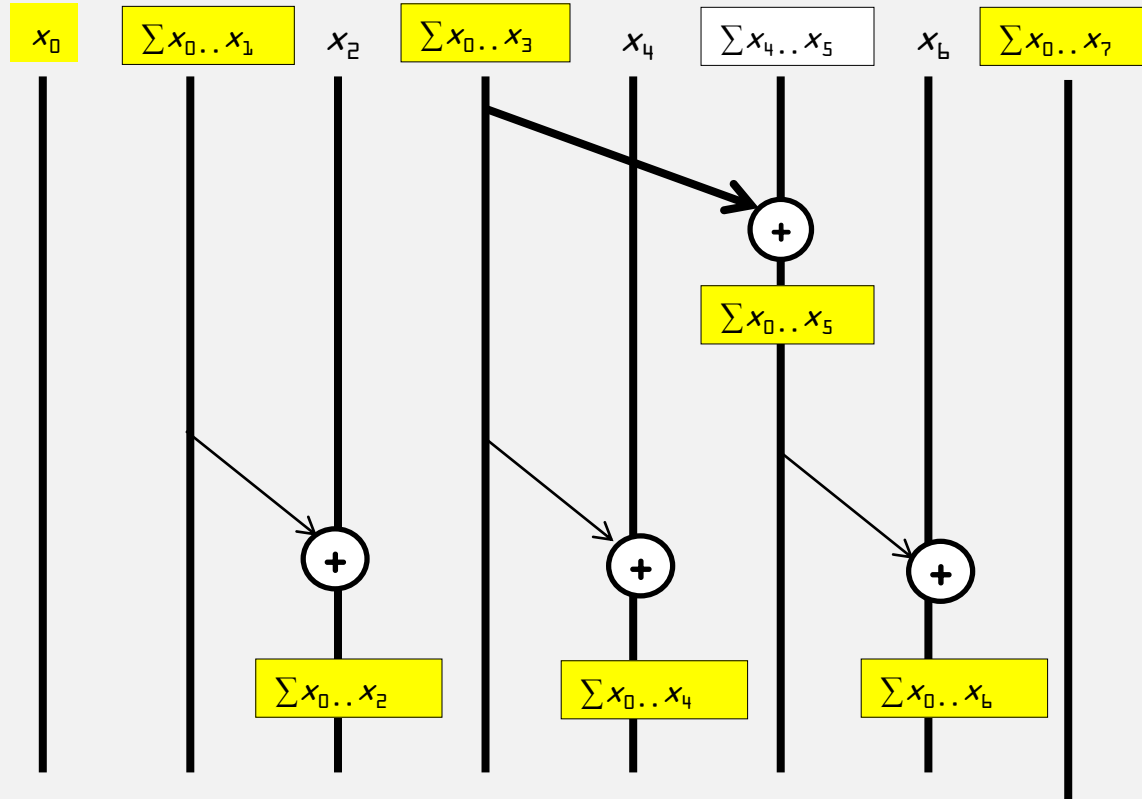
$\text{stride} = 1,$

$\text{index} = 1, 3, 5, 7, \dots$

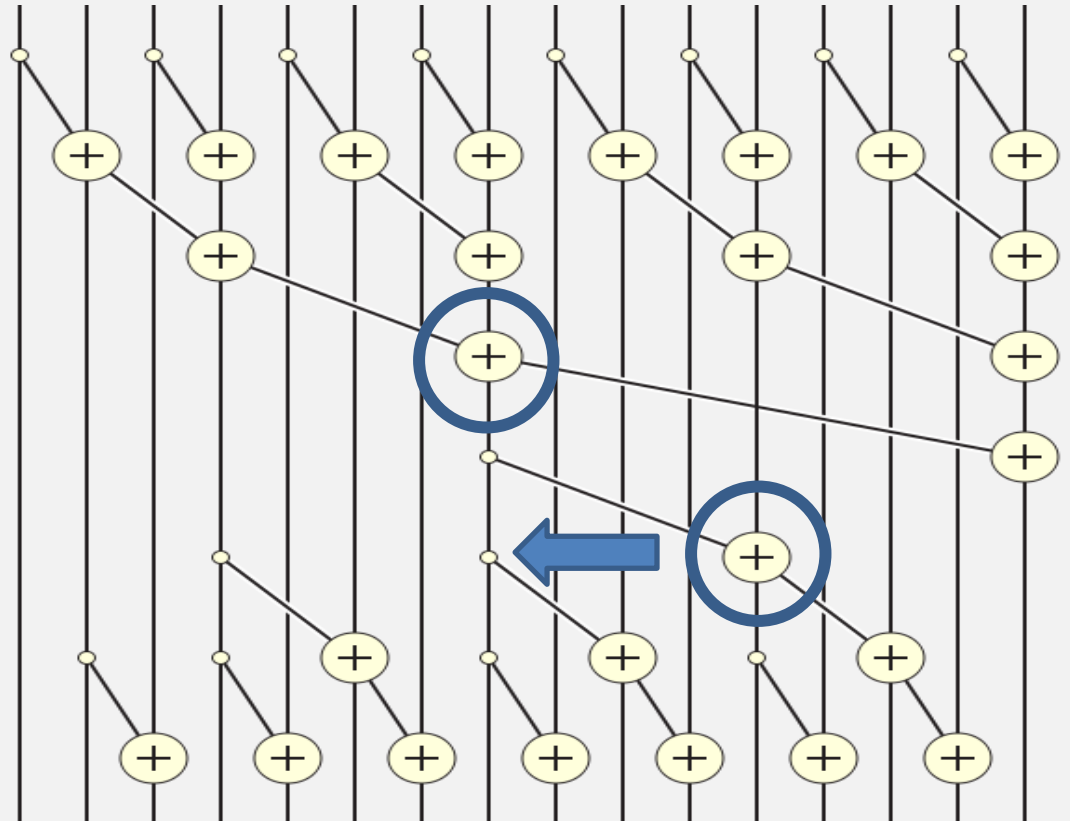
Parallel Scan - Post Reduction Reverse Phase



Parallel Scan - Post Reduction Reverse Phase



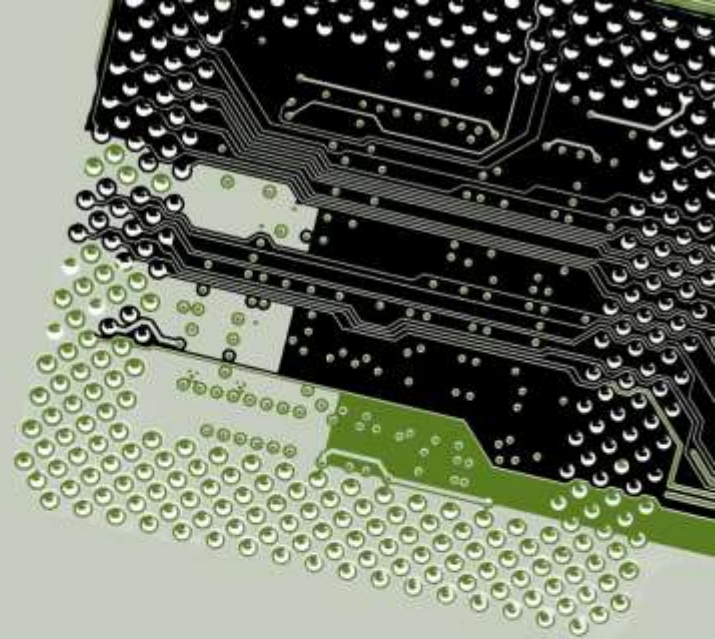
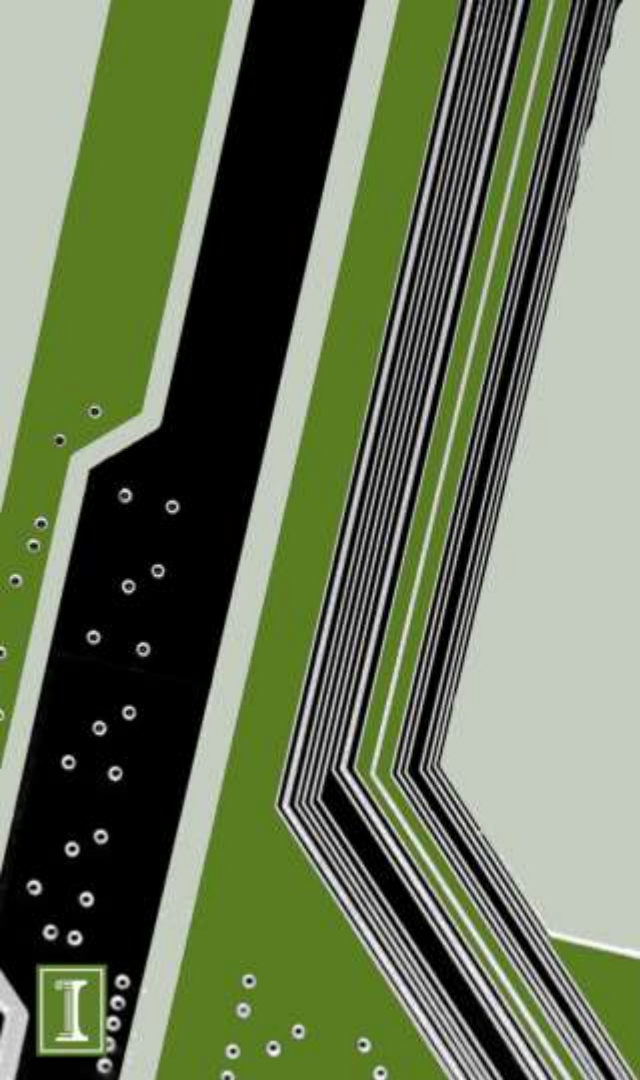
Putting it together



Post Reduction Reverse Phase Kernel Code

```
for (int stride = BLOCK_SIZE/2; stride > 0;
    stride /= 2) {
    __syncthreads();
    int index = (threadIdx.x+1)*stride*2 - 1;
    if(index+stride < 2*BLOCK_SIZE) {
        XY[index + stride] += XY[index];
    }
}
__syncthreads();
if (i < InputSize) Y[i] = XY[threadIdx.x];
```

First iteration for 16-element section
threadIdx.x = 0
stride = BLOCK_SIZE/2 = 8/2 = 4
index = 8-1 = 7



To learn more, read
Sections 9.4-9.5